

Name = Manoj Kumar

Batch = 1st September Batch

course = Data science Placement guarantee Course

Email = manojkumarrajp9990@gmail.com

```
In [195... import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
warnings.simplefilter(action='ignore', category=UserWarning)
```

Imported all the libraries which we need to use for this task

Part 1: Basic Data Exploration and Manipulation

```
In [2]: data = pd.read_csv("C:\\Users\\msgme\\Downloads\\Assignmnet_dataset.csv")
```

```
In [5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4000 entries, 0 to 3999
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   CustomerID            4000 non-null   int64
1   Name                  3340 non-null   object
2   Age                   3944 non-null   float64
3   Gender                2939 non-null   object
4   City                  3328 non-null   object
5   PurchaseAmount        3900 non-null   float64
6   PurchaseDate          4000 non-null   object
7   ProductCategory       3366 non-null   object
dtypes: float64(2), int64(1), object(5)
memory usage: 250.1+ KB
```

```
In [7]: data.head()
```

Out[7]:

	CustomerID	Name	Age	Gender	City	PurchaseAmount	PurchaseDate	ProductCat
0	8270	Jane Smith	46.0	Female	Los Angeles	648.27	2021-09-12 00:00:00	
1	1860	NaN	30.0	NaN	Houston	185.30	2020-06-30 00:00:00	Gro
2	6390	Jane Smith	38.0	Male	New York	564.92	2021-11-16 00:00:00	Home
3	6191	John Doe	75.0	Female	Houston	981.52	2021-11-11 00:00:00	
4	6734	NaN	38.0	NaN	Chicago	523.13	2020-09-04 00:00:00	Home

In [9]: data.tail()

Out[9]:

	CustomerID	Name	Age	Gender	City	PurchaseAmount	PurchaseDate	Prodi
3995	9347	Chris Johnson	64.0	NaN	Houston	NaN	2023-04-02 00:00:00	
3996	6592	John Doe	23.0	Female	New York	527.92	2023-09-20 00:00:00	
3997	2724	John Doe	61.0	Female	NaN	223.12	2020-11-10 00:00:00	
3998	4343	Emily White	25.0	NaN	Phoenix	721.52	2022-08-02 00:00:00	
3999	1886	John Doe	24.0	Female	Houston	93.93	2020-08-25 00:00:00	

In [11]: data.describe()

Out[11]:

	CustomerID	Age	PurchaseAmount
count	4000.000000	3944.000000	3900.000000
mean	5482.300250	47.774341	505.700428
std	2573.310642	18.834296	287.535960
min	1001.000000	-1.000000	10.250000
25%	3255.000000	32.000000	255.895000
50%	5492.500000	48.000000	506.750000
75%	7704.000000	64.000000	759.177500
max	9996.000000	79.000000	999.930000

1. Data Cleaning:

```
In [13]: missing_values= data.isnull().sum()
print(missing_values)
```


```
CustomerID      0
Name            660
Age             56
Gender          1061
City            672
PurchaseAmount  100
PurchaseDate     0
ProductCategory 634
dtype: int64
```

```
In [15]: print(missing_values[missing_values > 0])
```


```
Name            660
Age             56
Gender          1061
City            672
PurchaseAmount  100
ProductCategory 634
dtype: int64
```

```
In [17]: num_columns = ['CustomerID', 'Age', 'PurchaseAmount']
plt.figure(figsize=(10,8))
data[num_columns].hist(bins=20, figsize=(12, 8), layout=(3, 3))
plt.show()
```


<Figure size 1000x800 with 0 Axes>



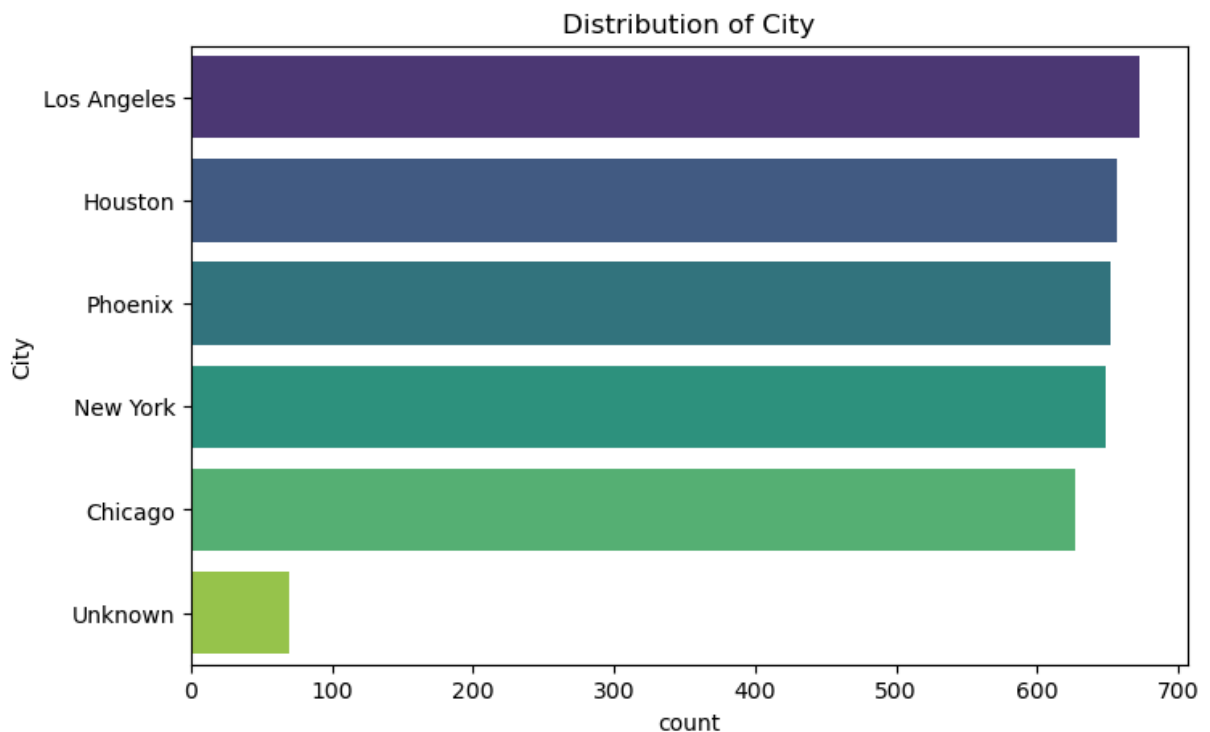
```
In [19]: plt.figure(figsize=(8, 5))
sns.heatmap(data[num_columns].corr(), annot=True, cmap='coolwarm', fmt='.2f')
plt.title("Correlation Heatmap")
plt.show()
```



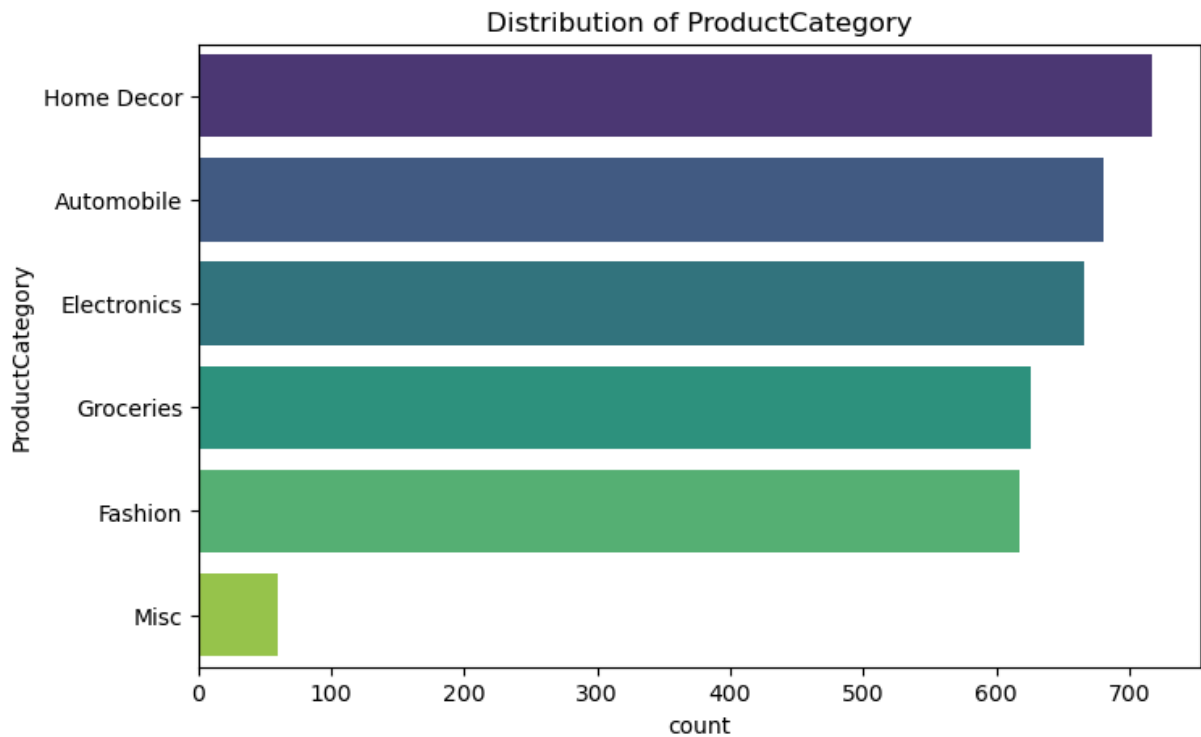
```
In [21]: property_ty = ['Gender']
for column in property_ty:
    plt.figure(figsize=(8, 5))
    sns.countplot(y=data[column], order=data[column].value_counts().index, palette='v')
    plt.title(f"Distribution of {column}")
    plt.show()
```



```
In [23]: property_ty = ['City']
for colmn in property_ty:
    plt.figure(figsize=(8, 5))
    sns.countplot(y=data[colmn], order=data[colmn].value_counts().index, palette='v
    plt.title(f"Distribution of {colmn}")
    plt.show()
```



```
In [25]: property_ty = ['ProductCategory']
for colmn in property_ty:
    plt.figure(figsize=(8, 5))
    sns.countplot(y=data[colmn], order=data[colmn].value_counts().index, palette='v'
    plt.title(f"Distribution of {colmn}")
    plt.show()
```



Handling missing values

Filling missing Name, Gender, City, and ProductCategory with 'Unknown'

```
In [29]: data['Name'].fillna('Unknown', inplace=True)
data['Gender'].fillna('Unknown', inplace=True)
data['City'].fillna('Unknown', inplace=True)
data['ProductCategory'].fillna('Unknown', inplace=True)
```

Filling missing Age with the median

```
In [33]: data['Age'].fillna(data['Age'].median(), inplace=True)
```

Filling missing PurchaseAmount with the mean

```
In [37]: data['PurchaseAmount'].fillna(data['PurchaseAmount'].mean(), inplace=True)
```

```
In [39]: missing_values = data.isnull().sum()
print(missing_values)
```

```

CustomerID      0
Name            0
Age            0
Gender         0
City           0
PurchaseAmount  0
PurchaseDate    0
ProductCategory 0
dtype: int64

```

Converting PurchaseDate to datetime format

```
In [43]: data['PurchaseDate'] = pd.to_datetime(data['PurchaseDate'], errors='coerce')
```

Keeping only the date and removing the time component

```
In [47]: data['PurchaseDate'] = data['PurchaseDate'].dt.date
```

```
In [49]: print(data['PurchaseDate'].head())
```

```

0    2021-09-12
1    2020-06-30
2    2021-11-16
3    2021-11-11
4    2020-09-04
Name: PurchaseDate, dtype: object

```

Removinhg duplicates from CustomerID and PurchaseDate columns

```
In [51]: data.drop_duplicates(subset=['CustomerID'], inplace=True)
```

```
In [53]: data.drop_duplicates(subset=['PurchaseDate'], inplace=True)
```

```
In [55]: data['CustomerID'].shape[0]
```

```
Out[55]: 1295
```

```
In [57]: data['PurchaseDate'].shape[0]
```

```
Out[57]: 1295
```

2: Customer Segmentation by City

```
In [61]: new_york_customers = data[data['City'] == 'New York']
```

```
In [63]: print("\nCustomers in New York:\n", new_york_customers[['CustomerID', 'Name', 'City']])
```

Customers in New York:

	CustomerID	Name	City
2	6390	Jane Smith	New York
6	1466	Jane Smith	New York
8	6578	Unknown	New York
21	9666	Jane Smith	New York
22	3558	Emily White	New York
...
3656	2812	Emily White	New York
3664	4340	John Doe	New York
3742	1991	Emily White	New York
3764	8876	Alex Brown	New York
3911	3153	Emily White	New York

[203 rows x 3 columns]

3: Total Purchase Amount

```
In [69]: total_purchase_amount = data['PurchaseAmount'].sum()
```

```
In [71]: print("\nTotal Purchase Amount:", total_purchase_amount)
```

Total Purchase Amount: 651161.903274359

4: Total Customer Count

```
In [75]: data['CustomerID'].sum()
```

Out[75]: 7148171

```
In [77]: total_unique_customers = data['CustomerID'].nunique()
```

```
In [79]: print("\nTotal Unique Customers:", total_unique_customers)
```

Total Unique Customers: 1295

5: Average Purchase Amount

```
In [83]: average_purchase_amount = data['PurchaseAmount'].mean()
```

```
In [85]: print("\nAverage Purchase Amount:", average_purchase_amount)
```

Average Purchase Amount: 502.8277245361845

Part 2: Intermediate Analysis and Aggregation

1. Purchase Analysis by Product Category

```
In [89]: total_purchase_by_category = data.groupby('ProductCategory')['PurchaseAmount'].sum()
         avg_age_by_category = data.groupby('ProductCategory')['Age'].mean()
```

```
In [91]: print("\nTotal Purchase by Product Category:\n", total_purchase_by_category)
         print("\nAverage Age by Product Category:\n", avg_age_by_category)
```


Total Purchase by Product Category:

```

ProductCategory
Automobile      117086.672569
Electronics     98657.482141
Fashion         103634.552141
Groceries       103681.531713
Home Decor     107606.572997
Misc           10470.680000
Unknown        110024.411713
Name: PurchaseAmount, dtype: float64

```

Average Age by Product Category:

```

ProductCategory
Automobile      47.839286
Electronics     47.690355
Fashion         47.251185
Groceries       48.046729
Home Decor     47.652174
Misc           52.681818
Unknown        47.736364
Name: Age, dtype: float64

```

2. Recent Purchase Analysis

```

In [95]: latest_date = data['PurchaseDate'].max()
         recent_purchases = data[data['PurchaseDate'] >= (latest_date - pd.Timedelta(days=30))

In [97]: print("\nRecent Purchases (Last 30 Days):\n", recent_purchases)

```

Recent Purchases (Last 30 Days):

	CustomerID	Name	Age	Gender	City	PurchaseAmount \
23	8849	Chris Johnson	28.0	Unknown	Unknown	915.55
77	3062	Emily White	43.0	Unknown	Houston	75.75
177	1202	Emily White	73.0	Unknown	Los Angeles	786.08
303	8806	Alex Brown	49.0	Unknown	Houston	436.41
386	8253	John Doe	24.0	Female	Los Angeles	935.73
543	6029	John Doe	68.0	Male	New York	19.46
569	5468	Jane Smith	60.0	Unknown	New York	494.38
588	4324	Unknown	58.0	Unknown	Houston	664.42
608	5784	Emily White	75.0	Unknown	New York	564.36
693	2887	Emily White	18.0	Male	Phoenix	172.01
791	9837	Emily White	24.0	Other	Phoenix	963.35
865	2970	Alex Brown	58.0	Unknown	Los Angeles	850.25
896	4756	Unknown	21.0	Unknown	Chicago	286.05
900	7898	Emily White	29.0	Female	Houston	918.46
912	1827	Chris Johnson	56.0	Unknown	Phoenix	407.90
1001	5873	Unknown	67.0	Male	Unknown	87.09
1104	9340	Emily White	55.0	Female	Houston	165.07
1461	8014	Unknown	18.0	Male	Los Angeles	886.43
1517	1158	Emily White	76.0	Male	Houston	473.61
2066	9706	Jane Smith	38.0	Male	Chicago	746.79
2392	1948	Jane Smith	36.0	Unknown	New York	44.65
2830	2588	Alex Brown	78.0	Other	Unknown	908.55
2912	9418	John Doe	23.0	Male	Unknown	489.58
2932	4220	Chris Johnson	45.0	Male	Phoenix	836.62
3030	8763	John Doe	72.0	Female	Los Angeles	184.31
3152	7696	Jane Smith	60.0	Other	Unknown	65.65
3217	2104	Emily White	36.0	Male	Los Angeles	924.24
3664	4340	John Doe	68.0	Unknown	New York	486.43

	PurchaseDate	ProductCategory
23	2023-12-22	Automobile
77	2023-12-09	Home Decor
177	2023-12-08	Unknown
303	2023-12-16	Fashion
386	2023-12-13	Electronics
543	2023-12-31	Automobile
569	2023-12-24	Automobile
588	2023-12-17	Misc
608	2023-12-26	Groceries
693	2023-12-28	Electronics
791	2023-12-04	Groceries
865	2023-12-05	Groceries
896	2023-12-23	Fashion
900	2023-12-30	Automobile
912	2023-12-10	Automobile
1001	2024-01-01	Unknown
1104	2023-12-27	Fashion
1461	2023-12-07	Unknown
1517	2023-12-19	Electronics
2066	2023-12-06	Fashion
2392	2023-12-29	Unknown
2830	2023-12-20	Automobile
2912	2023-12-14	Automobile
2932	2023-12-21	Groceries

3030	2023-12-15	Unknown
3152	2023-12-11	Unknown
3217	2023-12-12	Groceries
3664	2023-12-18	Home Decor

3. Gender-Based Purchase Analysis

```
In [101...] total_purchase_by_gender = data.groupby('Gender')['PurchaseAmount'].sum()
```

```
In [103...] print("\nTotal Purchase by Gender:\n", total_purchase_by_gender)
```

```
Total Purchase by Gender:
Gender
Female      153989.623426
Male        176721.202569
Other        156279.662997
Unknown     164171.414282
Name: PurchaseAmount, dtype: float64
```

4. Age-Based Purchase Segmentation

```
In [107...] def age_group(age):
    if age < 30:
        return 'Below 30'
    elif 30 <= age < 40:
        return '30-40'
    elif 40 <= age < 50:
        return '40-50'
    else:
        return '50+'
```

```
In [109...] data['AgeGroup'] = data['Age'].apply(age_group)
total_purchase_by_age_group = data.groupby('AgeGroup')['PurchaseAmount'].sum()
```

```
In [111...] print("\nTotal Purchase by Age Group:\n", total_purchase_by_age_group)
```

```
Total Purchase by Age Group:
AgeGroup
30-40      103970.981713
40-50      108606.292997
50+        305804.595995
Below 30    132780.032569
Name: PurchaseAmount, dtype: float64
```

5. Top Transactions

```
In [113...] top_transactions = data.nlargest(5, 'PurchaseAmount')[['CustomerID', 'Name', 'Product']]
```

```
In [115...] print("\nTop 5 Transactions:\n", top_transactions)
```

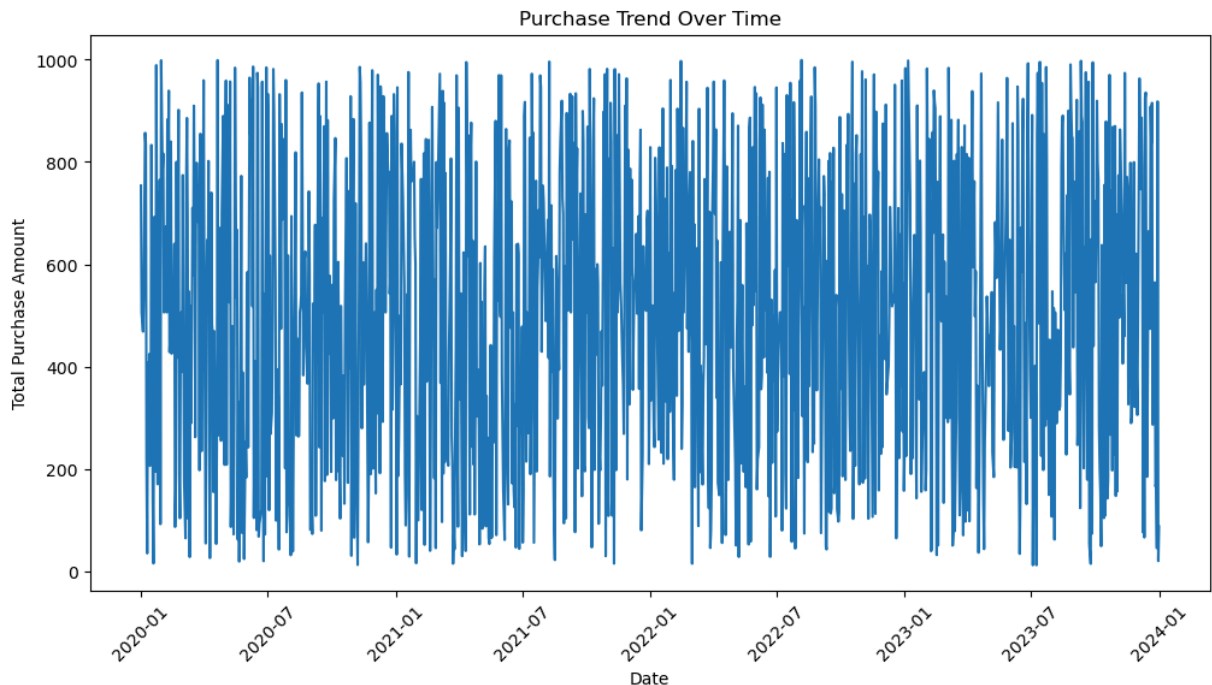
Top 5 Transactions:

	CustomerID	Name	ProductCategory	PurchaseAmount
2100	4489	Unknown	Electronics	999.46
3627	3890	John Doe	Misc	999.16
7	5426	Chris Johnson	Home Decor	998.86
1900	9648	Jane Smith	Groceries	998.29
905	1814	Emily White	Electronics	997.72

Part 3: Advanced Analysis and Insights

1. Purchase Trend Analysis

```
In [117... plt.figure(figsize=(12,6))
data.groupby('PurchaseDate')['PurchaseAmount'].sum().plot()
plt.xlabel('Date')
plt.ylabel('Total Purchase Amount')
plt.title('Purchase Trend Over Time')
plt.xticks(rotation=45)
plt.show()
```




2. City-Based Purchase Comparison

```
In [119... avg_purchase_by_city = data.groupby('City')['PurchaseAmount'].mean()
highest_avg_city = avg_purchase_by_city.idxmax()
```

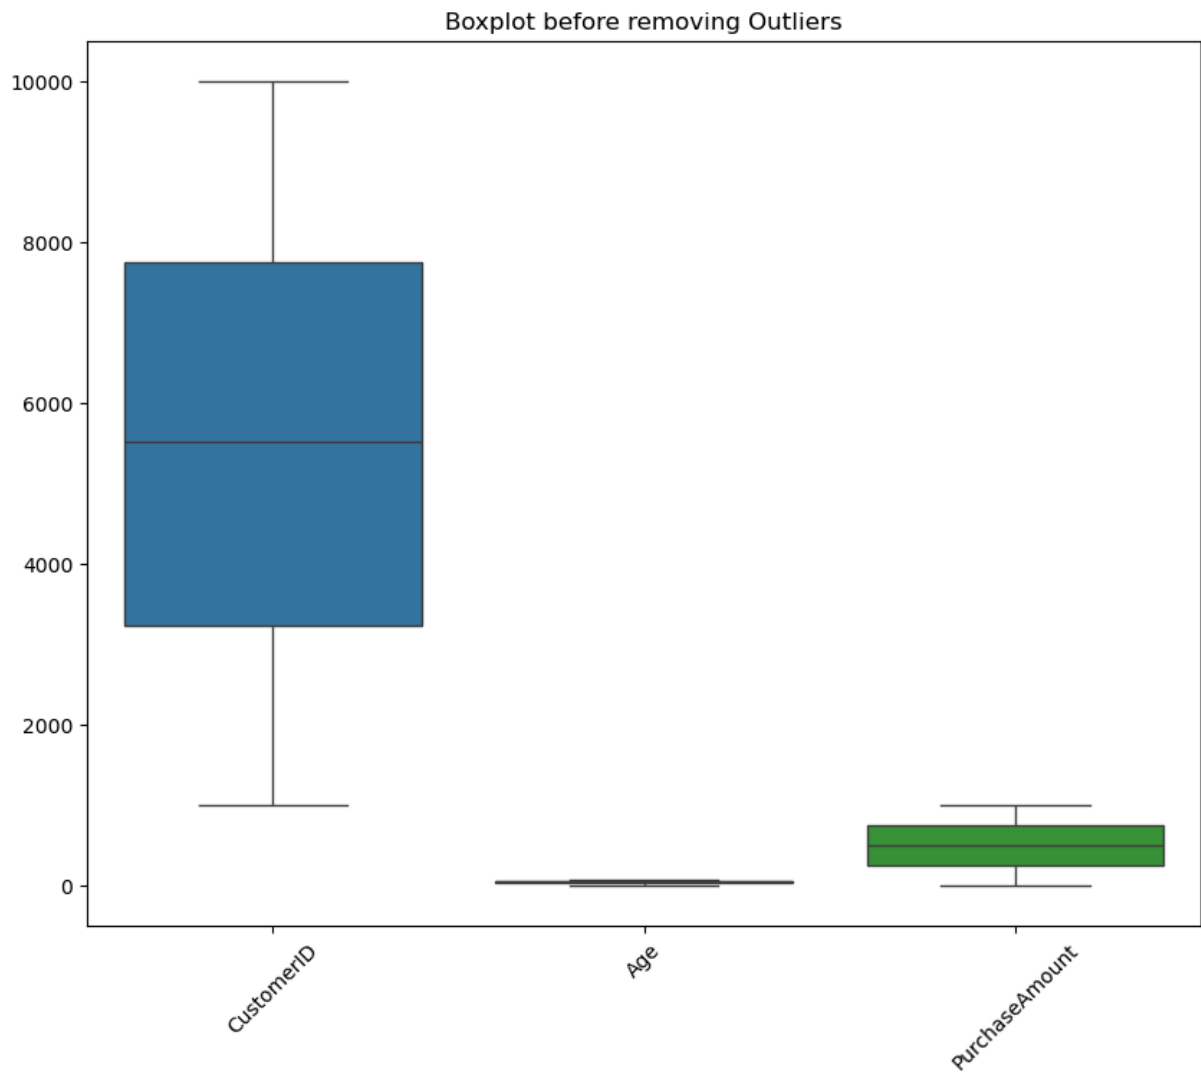
```
In [121... print("\nCity with Highest Average Purchase Amount:", highest_avg_city)
```

City with Highest Average Purchase Amount: Chicago

```
In [123... sns.histplot(data['PurchaseAmount'], kde=True)
plt.title('Histogram with Normal Distribution Curve')
plt.show()
```



```
In [125... plt.figure(figsize=(10, 8))
sns.boxplot(data=data[num_columns])
plt.title('Boxplot before removing Outliers')
plt.xticks(rotation=45)
plt.show()
```



as you can see there is no outliers in the data.

```
In [127...] data['ProductCategory']
```

```
Out[127...] 0      Unknown
            1      Groceries
            2      Home Decor
            3      Unknown
            4      Home Decor
            ...
            3925    Fashion
            3936    Groceries
            3938    Automobile
            3941    Fashion
            3948    Home Decor
            Name: ProductCategory, Length: 1295, dtype: object
```

4. Product Category Comparison (Statistical Analysis)

```
In [135...] import seaborn as sns
            from scipy import stats
```

Got the unique product categories

```
In [129... categories = data['ProductCategory'].unique()
```

```
In [131... print(categories)
```

```
['Unknown' 'Groceries' 'Home Decor' 'Fashion' 'Electronics' 'Automobile'
'Misc']
```

Created a list of data groups, excluding empty categories

And Checked if there are at least two groups to perform ANOVA

```
In [133... num_groups = len(categories)

print("Unique Product Categories:", categories)
print("Number of Groups:", num_groups)
```

```
Unique Product Categories: ['Unknown' 'Groceries' 'Home Decor' 'Fashion' 'Electronic
s' 'Automobile'
'Misc']
```

```
Number of Groups: 7
```

```
In [147... data_groups = [data[data['ProductCategory'] == cat]['PurchaseAmount'] for cat in ca
```

```
In [151... if len(data_groups) > 1:
    anova_result = stats.f_oneway(*data_groups)
    print("ANOVA Test Result:", anova_result)
else:
    print("Not enough data groups for ANOVA analysis")
```

```
ANOVA Test Result: F_onewayResult(statistic=0.5468369579646548, pvalue=0.77271933786
86309)
```

Got the result as "statistic=0.5468369579646548, pvalue=0.7727193378686309"

F- Statistic is 0.547 this value is the ratio in the groups and ratio between the groups.

This indicates that there can be significant difference between the groups .

P-Value is 0.773 the P-Value is a crucial measure used to access whether the results of this

test are statistically significant. It represents the likelihood that the observed outcomes occurred purely by chance.

If the P-Value is smaller than the chosen threshold (usually 0.05), we can reject the null hypothesis

and can say there is a statistically significant difference between the groups.

If P-Value is 0.05 or higher, we cannot reject null hypothesis and can say there is

no statistically significant difference between the groups.

P-Value is 0.773, which much higher than the usual significance level of 0.05, so we

fail to reject the null hypothesis, it means there is no strong evidence to suggest a

statistically significant difference in " PurchaseAmount" between the "ProductCategory" groups in the data

There is not a strong evidence that means PurchaseAmount are different across the product categories.

I am really sorry as did not have that time to visulize this in power BI as i got this project yesterday

evening so i did not have time to visulize it in power BI , if i will get time than forsure i can create

vizulization in Power BI as well.

In []:

In []: