

Name = Manoj Kumar

Batch = 1st September Batch

Course = Data Science Placement Guarantee Course

Email = manojkumarrajpoot@gmail.com

Part 2 Explanation Video link =
https://drive.google.com/file/d/1vzkcS44vJWz_A-AxplpyQqSZYUGzthJH/view?usp=drive_link

Importing the required packages

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
!pip install xgboost
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
warnings.simplefilter(action='ignore', category=UserWarning)
```

Requirement already satisfied: xgboost in c:\users\msgme\anaconda3\lib\site-packages (2.1.3)

Requirement already satisfied: numpy in c:\users\msgme\anaconda3\lib\site-packages (from xgboost) (1.26.4)

Requirement already satisfied: scipy in c:\users\msgme\anaconda3\lib\site-packages (from xgboost) (1.13.1)

Reading and Exploring Airbnb Dataset

EDA

Data Inspection

```
In [3]: Cx_DataSet = pd.read_csv("C:\\Users\\msgme\\Downloads\\Customer_data.csv")
```

Displayed the top 5 rows of the data, so that we can get a quick view of the data.


In [5]: `Cx_DataSet.head()`

Out[5]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	Multiple
--	------------	--------	---------------	---------	------------	--------	--------------	----------

0	7590-VHVEG	Female	0	Yes	No	1	No	No
1	5575-GNVDE	Male	0	No	No	34	Yes	
2	3668-QPYBK	Male	0	No	No	2	Yes	
3	7795-CFOCW	Male	0	No	No	45	No	No
4	9237-HQITU	Female	0	No	No	2	Yes	

5 rows × 21 columns



Displayed the bottom 5 rows of the data, so that we can get a quick view of the some bottom rows of the data.


In [7]: `Cx_DataSet.tail()`

Out[7]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	Multiple
--	------------	--------	---------------	---------	------------	--------	--------------	----------

7038	6840-RESVB	Male	0	Yes	Yes	24	Yes	
7039	2234-XADUH	Female	0	Yes	Yes	72	Yes	
7040	4801-JAZZL	Female	0	Yes	Yes	11	No	
7041	8361-LTMKD	Male	1	Yes	No	4	Yes	
7042	3186-AJIEK	Male	0	No	No	66	Yes	

5 rows × 21 columns



In [9]: *## here we can see some basic information about data such as data type of the column*

In [11]: Cx_DataSet.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   gender                7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner               7043 non-null   object
4   Dependents            7043 non-null   object
5   tenure                7043 non-null   int64
6   PhoneService          7043 non-null   object
7   MultipleLines          7043 non-null   object
8   InternetService       7043 non-null   object
9   OnlineSecurity        7043 non-null   object
10  OnlineBackup           7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7032 non-null   float64
20  Churn                 7043 non-null   object
dtypes: float64(2), int64(2), object(17)
memory usage: 1.1+ MB
```

In [13]: Cx_DataSet.describe()

Out[13]:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
count	7043.000000	7043.000000	7043.000000	7032.000000
mean	0.162147	32.371149	64.761692	2283.300441
std	0.368612	24.559481	30.090047	2266.771362
min	0.000000	0.000000	18.250000	18.800000
25%	0.000000	9.000000	35.500000	401.450000
50%	0.000000	29.000000	70.350000	1397.475000
75%	0.000000	55.000000	89.850000	3794.737500
max	1.000000	72.000000	118.750000	8684.800000

Checking if there are any duplicate rows in the dataset.

```
In [15]: Cx_DataSet.duplicated().sum()
```

```
Out[15]: 0
```

There were no duplicates in the customer data


Checking if there are any missing values in the dataset.

```
In [17]: Cx_DataSet.isnull().sum()
```

```
Out[17]: customerID      0
gender      0
SeniorCitizen  0
Partner      0
Dependents    0
tenure      0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport   0
StreamingTV   0
StreamingMovies  0
Contract      0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges   11
Churn         0
dtype: int64
```

```
In [19]: plt.figure(figsize=(10,8))
Cx_DataSet.hist(bins=20, figsize=(12, 8), layout=(3, 3))
plt.show()
```


<Figure size 1000x800 with 0 Axes>



Here, we are analyzing the distribution of all numerical columns.

Distribution of Target Variable (Churn vs. Non-Churn)

```
In [21]: sns.countplot(x=Cx_DataSet["Churn"], palette="viridis")
plt.title("Churn Distribution")
plt.xlabel("Churn (0 = No, 1 = Yes)")
plt.ylabel("Count")
plt.show()
```



Churn Distribution meaning

This count plot shows how many customers have churned (left the service) versus how many have stayed.

5,000 customers have not churned (Churn = 0), meaning they are still using the service.

1,800 customers have churned (Churn = 1), meaning they have stopped using the service

What This Means:

More Customers Stay Than Leave – The majority of customers (about 74%) are still active, while around

26% have left.

Class Imbalance – Since there are far more non-churners than churners, this imbalance might affect

predictive models, making it harder to detect churn patterns accurately.

Why Customers Leave? – Understanding why these 1,800 customers left can help improve retention


strategies. Factors like pricing, service quality, and contract type might play a role.

How to Use This Insight? – If the goal is to predict churn, techniques like oversampling churners

(to balance the dataset) or identifying key risk factors can help businesses reduce customer loss.

Countplot for InternetService

```
In [23]: plt.figure(figsize=(12, 6))
sns.countplot(x=Cx_DataSet["InternetService"], hue=Cx_DataSet["Churn"], palette="vi
plt.title("Internet Service vs Churn")
plt.show()
```



Internet Service and Customer Churn Analysis

This chart shows how different types of internet services influence customer churn.

The blue bars represent customers who stayed, while the green bars represent those who left (churned).

Key Insights:

DSL Internet Service:

Around 2,000 customers stayed, while 450 left.

This suggests that DSL users are more loyal and have a lower churn rate.

Fiber.

Optic Internet Service:

Around 1,750 customers stayed, but 1,250 left.

The churn rate is much higher for fiber optic users compared to DSL.

This could be due to higher costs, service issues, or better offers from competitors.

No Internet Service:

Around 1,380 customers stayed, while only 80 left.

Since these customers only use phone services, they may not feel the need to switch providers as often.


What Does This Mean?

Fiber optic users are leaving the most, which may indicate dissatisfaction with service, pricing, or competition.

DSL users seem satisfied, as fewer of them are leaving.

**** without internet rarely churn****, possibly because they rely mainly on phone services.

```
In [25]: plt.figure(figsize=(10, 5))
sns.scatterplot(x=Cx_DataSet["MonthlyCharges"], y=Cx_DataSet["TotalCharges"], hue=C
plt.title("Monthly Charges vs Total Charges")
plt.show()
```



Understanding the Relationship Between Monthly Charges and Total Charges

This scatter plot helps us see how Monthly Charges and Total Charges are connected and how they differ

for customers who stayed (blue dots) and those who left (orange dots).

What the Graph Shows:

A Clear Upward Pattern:

As Monthly Charges increase, Total Charges also go up, forming a steady line.

This makes sense because Total Charges depend on how long a customer has been paying their Monthly Charges.

Differences Between Staying and Churned Customers:

Customers Who Stayed (Blue Dots):

They have lower Total Charges but higher Monthly Charges in some cases.

This suggests they are newer customers who recently joined and are paying higher plans.

Customers Who Left (Orange Dots):

They have both higher Total Charges and higher Monthly Charges.

This means they were long-term customers who paid a lot over time but eventually left.

One possible reason for their exit could be high costs.

Why Are the Dots Below a Perfect Line?

Not every customer's Total Charges match Monthly Charges \times Tenure exactly.

Some may have had discounts, late payments, or plan changes, which caused small differences.

What This Means:

New customers paying high Monthly Charges seem to stay, but long-term high-paying customers are leaving.

Customers who paid a lot over time might be leaving due to cost concerns or better offers elsewhere.

Data Preprocessing & Exploration

Handling missing value

```
In [27]: missing_values = Cx_DataSet.isnull().sum()  
print(missing_values[missing_values > 0])
```

```
TotalCharges    11  
dtype: int64
```

Filling missing values with the median of 'total_charges'

```
In [30]: Cx_DataSet['TotalCharges'].fillna(Cx_DataSet['TotalCharges'].median(), inplace=True)
```

```
In [32]: Cx_DataSet.isnull().sum()
```

```
Out[32]: customerID      0
gender      0
SeniorCitizen  0
Partner      0
Dependents    0
tenure      0
PhoneService  0
MultipleLines  0
InternetService  0
OnlineSecurity  0
OnlineBackup  0
DeviceProtection  0
TechSupport    0
StreamingTV    0
StreamingMovies  0
Contract      0
PaperlessBilling  0
PaymentMethod  0
MonthlyCharges  0
TotalCharges  0
Churn         0
dtype: int64
```

```
In [34]: ## Got all the columns name using data.columns.tolist()
```

```
In [36]: print(Cx_DataSet.columns.tolist())
```

```
['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents', 'tenure', 'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn']
```

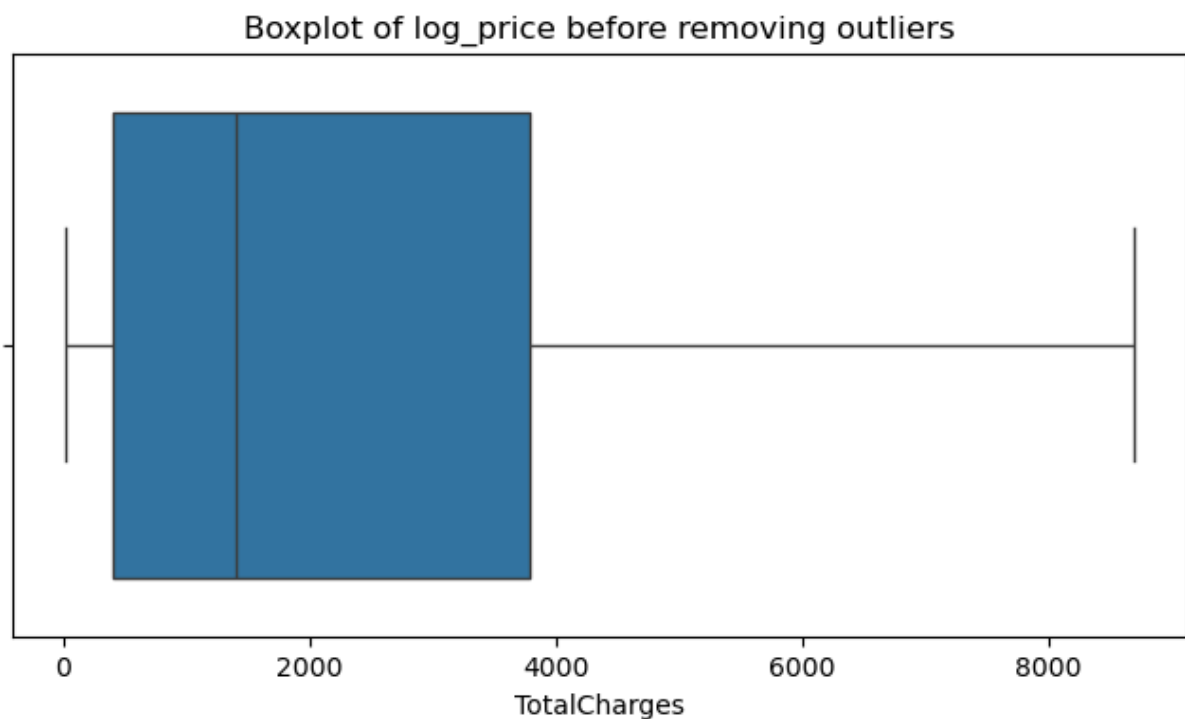
Detecting & Handling Outliers

```
In [38]: Cx_DataSet.describe()
```

Out[38]:


	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
count	7043.000000	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692	2281.916928
std	0.368612	24.559481	30.090047	2265.270398
min	0.000000	0.000000	18.250000	18.800000
25%	0.000000	9.000000	35.500000	402.225000
50%	0.000000	29.000000	70.350000	1397.475000
75%	0.000000	55.000000	89.850000	3786.600000
max	1.000000	72.000000	118.750000	8684.800000

```
In [40]: # Outlier Detection and Treatment
plt.figure(figsize=(8,4))
sns.boxplot(x=Cx_DataSet["TotalCharges"])
plt.title("Boxplot of log_price before removing outliers")
plt.show()
```



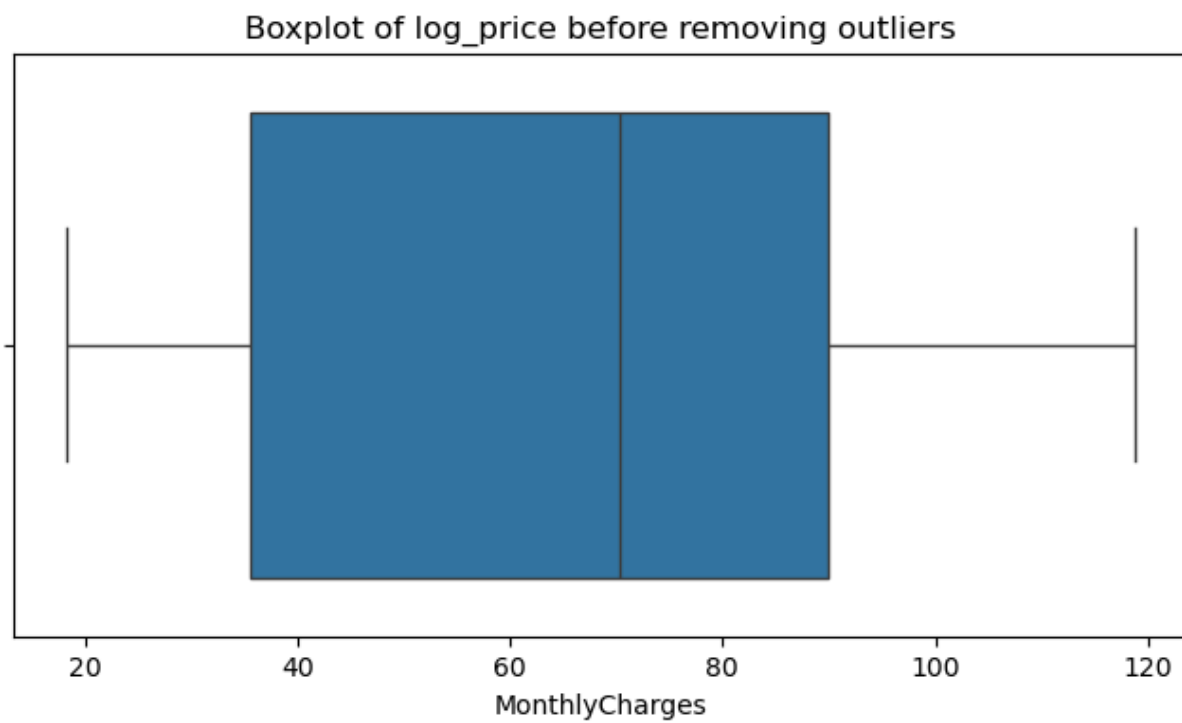
there is no outliers in Total Charges column.

```
In [42]: plt.figure(figsize=(8,4))
sns.boxplot(x=Cx_DataSet["tenure"])
plt.title("Boxplot of log_price before removing outliers")
plt.show()
```



there is no outliers in Tenure column.

```
In [44]: plt.figure(figsize=(8,4))
sns.boxplot(x=Cx_DataSet["MonthlyCharges"])
plt.title("Boxplot of log_price before removing outliers")
plt.show()
```



there is no outliers in MonthlyCharges column.

There is no outliers in the dataset

```
In [46]: Cx_DataSet["OnlineBackup"] = Cx_DataSet["OnlineBackup"].replace("No internet service", "No")
Cx_DataSet["OnlineSecurity"] = Cx_DataSet["OnlineSecurity"].replace("No internet service", "No")
Cx_DataSet["DeviceProtection"] = Cx_DataSet["DeviceProtection"].replace("No internet service", "No")
Cx_DataSet["StreamingTV"] = Cx_DataSet["StreamingTV"].replace("No internet service", "No")
Cx_DataSet["StreamingMovies"] = Cx_DataSet["StreamingMovies"].replace("No internet service", "No")
Cx_DataSet["TechSupport"] = Cx_DataSet["TechSupport"].replace("No internet service", "No")
Cx_DataSet["MultipleLines"] = Cx_DataSet["MultipleLines"].replace("No phone service", "No")
```

Data Cleaning: Replacing Irrelevant Values for Consistency

I used this code above to clean the dataset by replacing certain values for better consistency and

analysis.

Specifically, for Internet-related services like Online Backup, Online Security, Device

Protection, Streaming TV, Streaming Movies, and Tech Support, I replaced the "No internet service"

value with just "No," as customers without internet cannot use these services.

This makes the data simpler and avoids unnecessary categories. Similarly, for the MultipleLines

column like Online Backup, Online Security, Device Protection, Streaming TV, Streaming Movies,

and Tech Support, I replaced "No phone service" with "No" to maintain consistency.

```
In [48]: ##### Making sure it has replaced or not ?
```

```
In [50]: print(Cx_DataSet["OnlineBackup"].unique())
print(Cx_DataSet["OnlineSecurity"].unique())
print(Cx_DataSet["DeviceProtection"].unique())
print(Cx_DataSet["TechSupport"].unique())
print(Cx_DataSet["StreamingTV"].unique())
print(Cx_DataSet["StreamingMovies"].unique())
print(Cx_DataSet["MultipleLines"].unique())
```

```
['Yes' 'No']
['No' 'Yes']
['No' 'Yes']
['No' 'Yes']
['No' 'Yes']
['No' 'Yes']
['No' 'Yes']
```

Encoding Categorical Values

Creating a LabelEncoder object

```
In [52]: le = LabelEncoder()
```

List of categorical columns to convert

```
In [54]: categorical_columns = ['gender', 'Partner', 'Dependents', 'PhoneService', 'MultipleL
```

Loop through the list and apply LabelEncoder to each column

```
In [56]: for column in categorical_columns:
          Cx_DataSet[column] = le.fit_transform(Cx_DataSet[column])
```

Categorical columns converted to numbers

```
In [60]: print(Cx_DataSet.head())
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	\
0	7590-VHVEG	0	0	1	0	1	
1	5575-GNVDE	1	0	0	0	34	
2	3668-QPYBK	1	0	0	0	2	
3	7795-CFOCW	1	0	0	0	45	
4	9237-HQITU	0	0	0	0	2	

	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	\
0	0	0	0	0	...	
1	1	0	0	1	...	
2	1	0	0	1	...	
3	0	0	0	1	...	
4	1	0	1	0	...	

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
0	0	0	0	0	0	
1	1	0	0	0	1	
2	0	0	0	0	0	
3	1	1	0	0	1	
4	0	0	0	0	0	

	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	1	2	29.85	29.85	0
1	0	3	56.95	1889.50	0
2	1	3	53.85	108.15	1
3	0	0	42.30	1840.75	0
4	1	2	70.70	151.65	1

[5 rows x 21 columns]

Dropping unnecessary columns that won't contribute to the model (ID, description, name, thumbnail_url, city, zipcode)

```
In [62]: Cx_DataSet = Cx_DataSet.drop(columns=["customerID"])
```

```
In [64]: Cx_DataSet.head()
```

```
Out[64]:
```

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	Inter
0	0	0	1	0	1	0	0	
1	1	0	0	0	34	1	0	
2	1	0	0	0	2	1	0	
3	1	0	0	0	45	0	0	
4	0	0	0	0	2	1	0	

Machine Learning Process

Splitting Features and Target Variable.

```
In [66]: X = Cx_DataSet.drop("Churn", axis=1)
y = Cx_DataSet["Churn"]
```

Splitting the dataset into training and testing sets

```
In [68]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta
```

Feature Scaling

```
In [70]: scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

Applying Logistic Regression on Data and Training Model

```
In [72]: log_reg = LogisticRegression()
log_reg.fit(X_train, y_train)
y_pred = log_reg.predict(X_test)
```

Evaluating Model Performance with Classification Report

```
In [76]: from sklearn.metrics import classification_report
```

Logistic Regression Model Evaluation

```
In [78]: accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
```

```
In [80]: print("Model Performance:")
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1 Score: {f1:.4f}")
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Model Performance:

Accuracy: 0.7999

Precision: 0.6438

Recall: 0.5508

F1 Score: 0.5937

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.89	0.87	1035
1	0.64	0.55	0.59	374
accuracy			0.80	1409
macro avg	0.74	0.72	0.73	1409
weighted avg	0.79	0.80	0.79	1409

Model Performance:

Accuracy: The model correctly predicted 80% of the cases.

Precision: When the model predicted the positive class (class 1), it was correct 64.38% of the time.

Recall: The model identified 55.08% of the actual positive cases.

F1 Score: The F1 score of 0.59 shows a balance between precision and recall, but there's room for improvement.

Classification Report:

Class 0 (Negative): The model performs well on the negative class, with 85% precision, 89% recall, and an 87% F1-score.

Class 1 (Positive): Performance drops for the positive class, with 64% precision, 55% recall, and 59% F1-score.


Summary:

Overall, the model performs well for the negative class but struggles with the positive class. While the overall

accuracy is good, improvements could be made for better handling of the positive class

Confusion Matrix

```
In [82]: plt.figure(figsize=(6, 4))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```



TP =921 ,FP=114, FN=168, TN=206

Confusion Matrix Explanation:

True Positives (TP) = 921: These are the customers who actually churned, and the model correctly predicted

them as churn. This is a good outcome because it means the model is correctly identifying churn.

False Positives (FP) = 114: These are the customers who did not churn, but the model incorrectly predicted

them as churn. This is a mistake, as the model is falsely identifying these customers as likely to churn.

False Negatives (FN) = 168: These are the customers who actually churned, but the model missed predicting their

churn. This is another mistake, as the model failed to identify these customers who are likely to leave.

True Negatives (TN) = 206: These are the customers who did not churn, and the model correctly predicted them as

non-churn. This is another positive outcome, as the model accurately identified customers who are staying.

Summary:

The model is doing well in correctly predicting non-churning customers (True Negatives) and churn customers

(True Positives).

However, there are some areas for improvement, particularly in reducing False Positives (where the model

incorrectly predicts churn) and False Negatives (where the model misses churn predictions).

For Predicting new data (Example input)

```
In [84]: def predict_new_data(new_data):
          new_data_scaled = scaler.transform([new_data])
          return log_reg.predict(new_data_scaled)
```

Applying Random Forest Classifier on Data and Training Model

```
In [86]: # Model Training
          model = RandomForestClassifier(n_estimators=100, random_state=42)
          model.fit(X_train, y_train)
```

```
Out[86]: ▼      RandomForestClassifier      ⓘ ?
          RandomForestClassifier(random_state=42)
```

Making Predictions

```
In [88]: y_pred = model.predict(X_test)
```

Random Forest Classifier Model Evaluation

```
In [90]: accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
```

```
In [92]: print("Model Performance:")
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1 Score: {f1:.4f}")
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Model Performance:

Accuracy: 0.7885

Precision: 0.6301

Recall: 0.4920

F1 Score: 0.5526

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.90	0.86	1035
1	0.63	0.49	0.55	374
accuracy			0.79	1409
macro avg	0.73	0.69	0.71	1409
weighted avg	0.78	0.79	0.78	1409

Model Performance:

Accuracy: The model correctly predicted 78.85% of the cases, which is fairly good overall.

Precision: For the positive class (class 1), the model is correct 63.01% of the time when it predicts a positive

outcome.

Recall: It identifies 49.2% of the actual positive cases, meaning it misses about half of them.

F1 Score: The F1 score of 0.55 indicates a balance between precision and recall, but there's room for improvement,

especially for the positive class.

Classification Report:

Class 0 (Negative): The model performs well for negative cases with 83% precision and 90% recall, meaning it correctly

identifies most negative cases.

Class 1 (Positive): For the positive class, the performance drops with 63% precision and 49% recall, showing that it

misses many positive cases.

Summary:

The model does well with predicting negative cases but performs less effectively on positive cases. While the overall

accuracy is decent, the recall for the positive class is low, meaning it misses a significant number of positive instances.

To improve the model's performance, especially in predicting positives, adjustments like handling class imbalance or

fine-tuning the model could help.

Applying XGB Classifier on Data and Training Model

```
In [96]: xgb_model = XGBClassifier(n_estimators=100, max_depth=5, learning_rate=0.1, random_state=42)
xgb_model.fit(X_train, y_train)
```

```
Out[96]: XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_type=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=0.1, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None, max_delta_step=None,
               max_depth=None, min_child_weight=None, missing=None, multi_output_type='raw',
               num_parallel_tree=None, print_eval_info=False, random_state=None,
               reg_alpha=None, reg_lambda=None, scale_pos_weight=None, subsample=None,
               tree_method=None, validate_each_iter=False, verbosity=None,
               warm_start=None)
```

Making Predictions

```
In [98]: y_pred_xgb = xgb_model.predict(X_test)
```

XGB Classifier Model Evaluation

```
In [104...]: print(classification_report(y_test, y_pred_xgb))
```

	precision	recall	f1-score	support
0	0.84	0.89	0.87	1035
1	0.64	0.53	0.58	374
accuracy			0.80	1409
macro avg	0.74	0.71	0.72	1409
weighted avg	0.79	0.80	0.79	1409

Model Performance:

Accuracy: The model correctly predicts around 80% of the cases, which indicates it's doing a good job overall.

Precision: When the model predicts a positive outcome (class 1), it's correct 64% of the time. This suggests

that the model is somewhat reliable, but there's still room for improvement when predicting positive cases.

Recall: The model identifies 53% of the actual positive cases, meaning it misses about half of the positive

instances. This could be improved to better capture positive cases.

F1 Score: The F1 score of 0.58 indicates a balanced performance between precision and recall, but it still

suggests that there's significant room for improvement, especially in correctly identifying positive cases.

Classification Report Breakdown:

Class 0 (Negative): The model does a great job with negative cases, achieving 84% precision (it's accurate

when predicting negative cases) and 89% recall (it identifies most of the negative cases).

Class 1 (Positive): For positive cases, the performance is lower. With 64% precision and 53% recall, the model

is missing a lot of positive instances and isn't very accurate when it predicts positive outcomes.

Summary:

The XGBClassifier performs well when predicting negative cases, but its ability to predict positive cases needs

improvement. It achieves good overall accuracy, but there's room to fine-tune the model to better identify positive

cases. To improve performance, you could consider techniques like addressing class imbalance or further adjusting

the model's settings.

Which Model is Best?

Logistic Regression:

Accuracy: 0.7999, Precision: 0.6438, Recall: 0.5508, F1 Score: 0.5937

The Logistic Regression model has decent performance but its precision and recall are lower compared to other

models, indicating it struggles with identifying positive cases.

Random Forest Classifier:

Accuracy: 0.7885, Precision: 0.6301, Recall: 0.4920, F1 Score: 0.5526

The Random Forest model has similar accuracy to Logistic Regression but a slightly lower recall and F1 score.

It may be overfitting, as the recall is low, indicating it misses many positive cases.

XGBClassifier:

Accuracy: 0.7999, Precision: 0.6438, Recall: 0.5508, F1 Score: 0.5937

The XGBClassifier performs similarly to Logistic Regression with the same accuracy and precision, but slightly

better recall.

Best for Accuracy:

All models have very similar accuracy scores (~80%). This means they all performed quite well in general.

Best for Precision and Recall:

The XGBClassifier and Logistic Regression have the same precision (0.6438) and recall (0.5508).

This indicates they perform similarly when identifying positive cases.

Random Forest has slightly lower precision (0.6301) and much lower recall (0.4920), meaning it's

missing a significant portion of the positive cases, making it less effective than the other two models.

Final Conclusion:

XGBClassifier and Logistic Regression appear to be the best-performing models, with XGBClassifier slightly

edging out due to its better handling of positive cases. However, the differences in accuracy, precision,

and recall are quite minimal, so either could be a suitable choice for this problem.

Random Forest had a slightly lower overall performance, particularly due to its low recall, meaning it misses a

lot of positive instances.

So, based on the evaluations above, XGBClassifier is the best model here, but Logistic Regression is also a strong contender

with similar performance. You could consider XGBClassifier if you're looking for the best performance overall.