

Deep Learning Assignment : Image Classification Using CNNs

Name = Manoj Kumar

Batch = 1st September Batch

Course = Data Science Placement Guarantee Course

Email = manojkumarrajput@gmail.com

Objective

This assignment aims to help you understand the basics of Convolutional Neural Networks

(CNNs), their implementation, and evaluation in image classification tasks

Instructions

Complete the tasks below. Each task specifies the marks assigned. Submit your code,

outputs, and a brief explanation for each step. Use the CIFAR-10 dataset for this assignment.

Importing the required packages

```
In [1]: import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt
import numpy as np
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.regularizers import l2
from sklearn.model_selection import train_test_split
from tensorflow.keras.callbacks import EarlyStopping
```

Loading the dataset

```
In [3]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

Task 1: Data Exploration and Preparation

Verifying the shapes

```
In [7]: print("Training data shape:", x_train.shape)
print("Test data shape:", x_test.shape)
```

Training data shape: (50000, 32, 32, 3)
 Test data shape: (10000, 32, 32, 3)

Count of unique labels.

```
In [10]: print("Unique labels:", np.unique(y_train))
```

Unique labels: [0 1 2 3 4 5 6 7 8 9]

Displaying 5 sample images with labels

```
In [13]: def show_images(x_train, y_train, x_test, y_test):
    class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                   'dog', 'frog', 'horse', 'ship', 'truck']

    plt.figure(figsize=(10, 5))
    for i in range(5):
        plt.subplot(1, 5, i+1)
        plt.imshow(x_train[i])
        plt.title(class_names[y_train[i][0]])
        plt.axis('off')

    plt.show()

    return x_train, y_train, x_test, y_test
```

displaying sample images using "show_image" function which i created above from the training and testing data to

help visualize the dataset.

```
In [16]: x_train, y_train, x_test, y_test = show_images(x_train, y_train, x_test, y_test)
```



Printing the count of unique labels

```
In [19]: unique_labels, counts = np.unique(y_train, return_counts=True)
print("Unique Labels:", unique_labels)
print("Label Distribution:", counts)

Unique Labels: [0 1 2 3 4 5 6 7 8 9]
Label Distribution: [5000 5000 5000 5000 5000 5000 5000 5000 5000 5000]
```

Preprocess Data

1. Normalizing pixel values to [0,1] range

Here i am making a function and this function prepares the data before training. It first scales the pixel values between 0 and 1

to help the model learn better. Then, it splits 5000 images from the training set for validation, so we can check the model's progress.

The rest is used for training, while the test set stays the same for final evaluation. This process helps improve accuracy and performance of the model.

```
In [24]: def preprocess_data(x_train, y_train, x_test, y_test):
    x_train, x_test = x_train / 255.0, x_test / 255.0
    x_val, y_val = x_train[:5000], y_train[:5000]
    x_train, y_train = x_train[5000:], y_train[5000:]

    return x_train, y_train, x_val, y_val, x_test, y_test
```

Converting data to float and dividing by 255 to ensures all values are between 0 and 1, making it easier for the model to process.

```
In [27]: def normalize_data(x):
    if x.max() <= 1.0:
        print("Data is already normalized! Skipping normalization.")
        return x
    x = x.astype('float32') / 255.0
    return x
```

```
In [29]: x_train = normalize_data(x_train)
x_test = normalize_data(x_test)
```

Splitting training data into 80% training and 20% validation

```
In [32]: def split_train_validation(x_train, y_train, test_size=0.2, random_state=42):
    x_train_new, x_val, y_train_new, y_val = train_test_split(x_train, y_train, tes
    return x_train_new, y_train_new, x_val, y_val
```

```
In [34]: x_train, y_train, x_val, y_val = split_train_validation(x_train, y_train)
```

```
In [36]: x_train.min()
```

```
Out[36]: 0.0
```

```
In [38]: x_train.max()
```

```
Out[38]: 1.0
```

```
In [40]: print("Training Set Shape:", x_train.shape, y_train.shape)
print("Validation Set Shape:", x_val.shape, y_val.shape)
```

Training Set Shape: (40000, 32, 32, 3) (40000, 1)

Validation Set Shape: (10000, 32, 32, 3) (10000, 1)

Task 2: Build and Train a CNN Model

```
In [43]: def build_model():
    model = models.Sequential([
        layers.Conv2D(64, (3,3), activation='relu', padding='same', input_shape=(32,
            kernel_regularizer=l2(0.0005)),
        layers.BatchNormalization(),
        layers.MaxPooling2D((2,2)),
        layers.Dropout(0.3),

        layers.Conv2D(128, (3,3), activation='relu', padding='same', kernel_regularizer=
        layers.BatchNormalization(),
        layers.MaxPooling2D((2,2)),
        layers.Dropout(0.3),

        layers.Conv2D(256, (3,3), activation='relu', padding='same', kernel_regularizer=
        layers.BatchNormalization(),
        layers.MaxPooling2D((2,2)),
        layers.Dropout(0.4),

        layers.Flatten(),
        layers.Dense(512, activation='relu', kernel_regularizer=l2(0.0005)),
        layers.Dropout(0.4),
        layers.Dense(10, activation='softmax')])
    model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics
    return model
```

```
In [45]: model = build_model()
model.summary()
```

```
C:\Users\msgme\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base_conv.
py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. W
hen using Sequential models, prefer using an `Input(shape)` object as the first laye
r in the model instead.
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "sequential"
```

Layer (type)	Output Shape
conv2d (Conv2D)	(None, 32, 32, 64)
batch_normalization (BatchNormalization)	(None, 32, 32, 64)
max_pooling2d (MaxPooling2D)	(None, 16, 16, 64)
dropout (Dropout)	(None, 16, 16, 64)
conv2d_1 (Conv2D)	(None, 16, 16, 128)
batch_normalization_1 (BatchNormalization)	(None, 16, 16, 128)
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 128)
dropout_1 (Dropout)	(None, 8, 8, 128)
conv2d_2 (Conv2D)	(None, 8, 8, 256)
batch_normalization_2 (BatchNormalization)	(None, 8, 8, 256)
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 256)
dropout_2 (Dropout)	(None, 4, 4, 256)
flatten (Flatten)	(None, 4096)
dense (Dense)	(None, 512)
dropout_3 (Dropout)	(None, 512)
dense_1 (Dense)	(None, 10)



Total params: 2,475,402 (9.44 MB)

Trainable params: 2,474,506 (9.44 MB)

Non-trainable params: 896 (3.50 KB)

CNN Model Summary Explanation:

This model is built to classify images from the CIFAR-10 dataset. It has three layers that detect patterns in images,

like edges and shapes. Batch Normalization helps the model learn faster, and MaxPooling reduces image size to focus on

important details. Dropout layers prevent overfitting, so the model doesn't just memorize but learns properly. The final

part of the model flattens the data and passes it through a dense layer before making a prediction. Overall, it has about

2.4 million trainable settings, making it strong enough to recognize different objects

Training the Model

Plotting training & validation loss and accuracy

```
In [49]: early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights
```

```
In [51]: def train_model(model, x_train, y_train, x_val, y_val):
    history = model.fit(x_train, y_train, epochs=20, validation_data=(x_val, y_val))

    plt.figure(figsize=(12, 4))

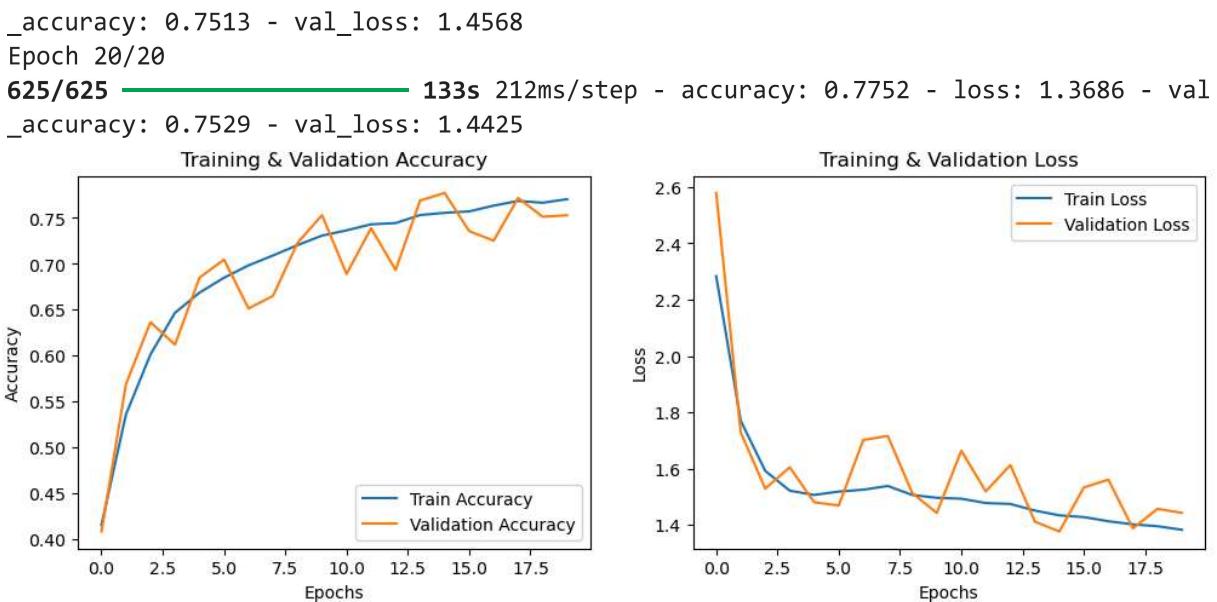
    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'], label='Train Accuracy')
    plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')
    plt.legend()
    plt.title('Training & Validation Accuracy')

    plt.subplot(1, 2, 2)
    plt.plot(history.history['loss'], label='Train Loss')
    plt.plot(history.history['val_loss'], label='Validation Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.legend()
    plt.title('Training & Validation Loss')

    plt.show()
```

```
In [53]: train_model(model, x_train, y_train, x_val, y_val)
```

Epoch 1/20
625/625 165s 248ms/step - accuracy: 0.3439 - loss: 2.8669 - val_accuracy: 0.4084 - val_loss: 2.5792
Epoch 2/20
625/625 155s 247ms/step - accuracy: 0.5193 - loss: 1.8250 - val_accuracy: 0.5690 - val_loss: 1.7282
Epoch 3/20
625/625 154s 247ms/step - accuracy: 0.5946 - loss: 1.6031 - val_accuracy: 0.6365 - val_loss: 1.5286
Epoch 4/20
625/625 154s 246ms/step - accuracy: 0.6443 - loss: 1.5116 - val_accuracy: 0.6120 - val_loss: 1.6044
Epoch 5/20
625/625 154s 246ms/step - accuracy: 0.6668 - loss: 1.4796 - val_accuracy: 0.6849 - val_loss: 1.4801
Epoch 6/20
625/625 154s 246ms/step - accuracy: 0.6830 - loss: 1.5075 - val_accuracy: 0.7047 - val_loss: 1.4688
Epoch 7/20
625/625 151s 241ms/step - accuracy: 0.7012 - loss: 1.5071 - val_accuracy: 0.6514 - val_loss: 1.7014
Epoch 8/20
625/625 124s 198ms/step - accuracy: 0.7143 - loss: 1.5174 - val_accuracy: 0.6651 - val_loss: 1.7158
Epoch 9/20
625/625 128s 204ms/step - accuracy: 0.7284 - loss: 1.4817 - val_accuracy: 0.7229 - val_loss: 1.5128
Epoch 10/20
625/625 129s 206ms/step - accuracy: 0.7370 - loss: 1.4715 - val_accuracy: 0.7531 - val_loss: 1.4423
Epoch 11/20
625/625 130s 209ms/step - accuracy: 0.7398 - loss: 1.4779 - val_accuracy: 0.6890 - val_loss: 1.6633
Epoch 12/20
625/625 130s 208ms/step - accuracy: 0.7498 - loss: 1.4550 - val_accuracy: 0.7389 - val_loss: 1.5187
Epoch 13/20
625/625 134s 214ms/step - accuracy: 0.7569 - loss: 1.4426 - val_accuracy: 0.6933 - val_loss: 1.6126
Epoch 14/20
625/625 132s 211ms/step - accuracy: 0.7607 - loss: 1.4308 - val_accuracy: 0.7689 - val_loss: 1.4117
Epoch 15/20
625/625 132s 212ms/step - accuracy: 0.7601 - loss: 1.4184 - val_accuracy: 0.7773 - val_loss: 1.3760
Epoch 16/20
625/625 133s 212ms/step - accuracy: 0.7628 - loss: 1.4043 - val_accuracy: 0.7357 - val_loss: 1.5329
Epoch 17/20
625/625 132s 211ms/step - accuracy: 0.7727 - loss: 1.3868 - val_accuracy: 0.7253 - val_loss: 1.5605
Epoch 18/20
625/625 132s 212ms/step - accuracy: 0.7754 - loss: 1.3752 - val_accuracy: 0.7718 - val_loss: 1.3873
Epoch 19/20
625/625 133s 213ms/step - accuracy: 0.7738 - loss: 1.3743 - val_accuracy: 0.7700 - val_loss: 1.3860



The model started with low accuracy (34.39%) but kept improving, reaching 77.52% on training data and 75.29% on validation data.

The loss went down, meaning the model learned well, but the ups and downs in validation accuracy show some overfitting.

To make it better, I have tried techniques like dropout, adjusting the learning rate, or adding regularization.

Creating a Function to Detect Overfitting by Comparing Training and Validation Accuracy

Creating a **function** below which checks if the model is overfitting by comparing how well it performs on training vs.

validation data.

It makes predictions on both training and validation sets.

Converts the predictions into actual class labels.

Compares predicted labels with real labels to find accuracy.

Prints both accuracies to see how well the model is learning.

If training accuracy is much higher than validation accuracy, the model might be overfitting, it means memorizing training

data instead of learning patterns.

In [56]:

```
def check_overfitting(model, x_train, y_train, x_val, y_val):
    train_acc = np.mean(np.argmax(model.predict(x_train), axis=1) == y_train.flatten())
    val_acc = np.mean(np.argmax(model.predict(x_val), axis=1) == y_val.flatten())
```

```

    print(f"Train Accuracy: {train_acc:.4f}")
    print(f"Validation Accuracy: {val_acc:.4f}")

    return train_acc, val_acc

```

Adding Comment briefly on overfitting or underfitting based on the plots

```
In [59]: def analyze_performance(train_acc, val_acc):
    if train_acc > val_acc + 0.1:
        print(" Possible Overfitting: Training accuracy is much higher than validation accuracy")
    elif val_acc > train_acc:
        print(" Possible Underfitting: Validation accuracy is higher than training accuracy")
    else:
        print(" Model is generalizing well without significant overfitting or underfitting")
```

```
In [61]: train_acc, val_acc = check_overfitting(model, x_train, y_train, x_val, y_val)
analyze_performance(train_acc, val_acc)
```

```
1250/1250 ————— 24s 19ms/step
313/313 ————— 7s 22ms/step
Train Accuracy: 0.8409
Validation Accuracy: 0.7773
Model is generalizing well without significant overfitting or underfitting.
```

Task 3: Evaluate the Model

```
In [64]: def evaluate_model(model, x_test, y_test):
    test_loss, test_acc = model.evaluate(x_test, y_test)
    print("Test Accuracy:", test_acc)
```

Generating Confusion Matrix and Classification Report

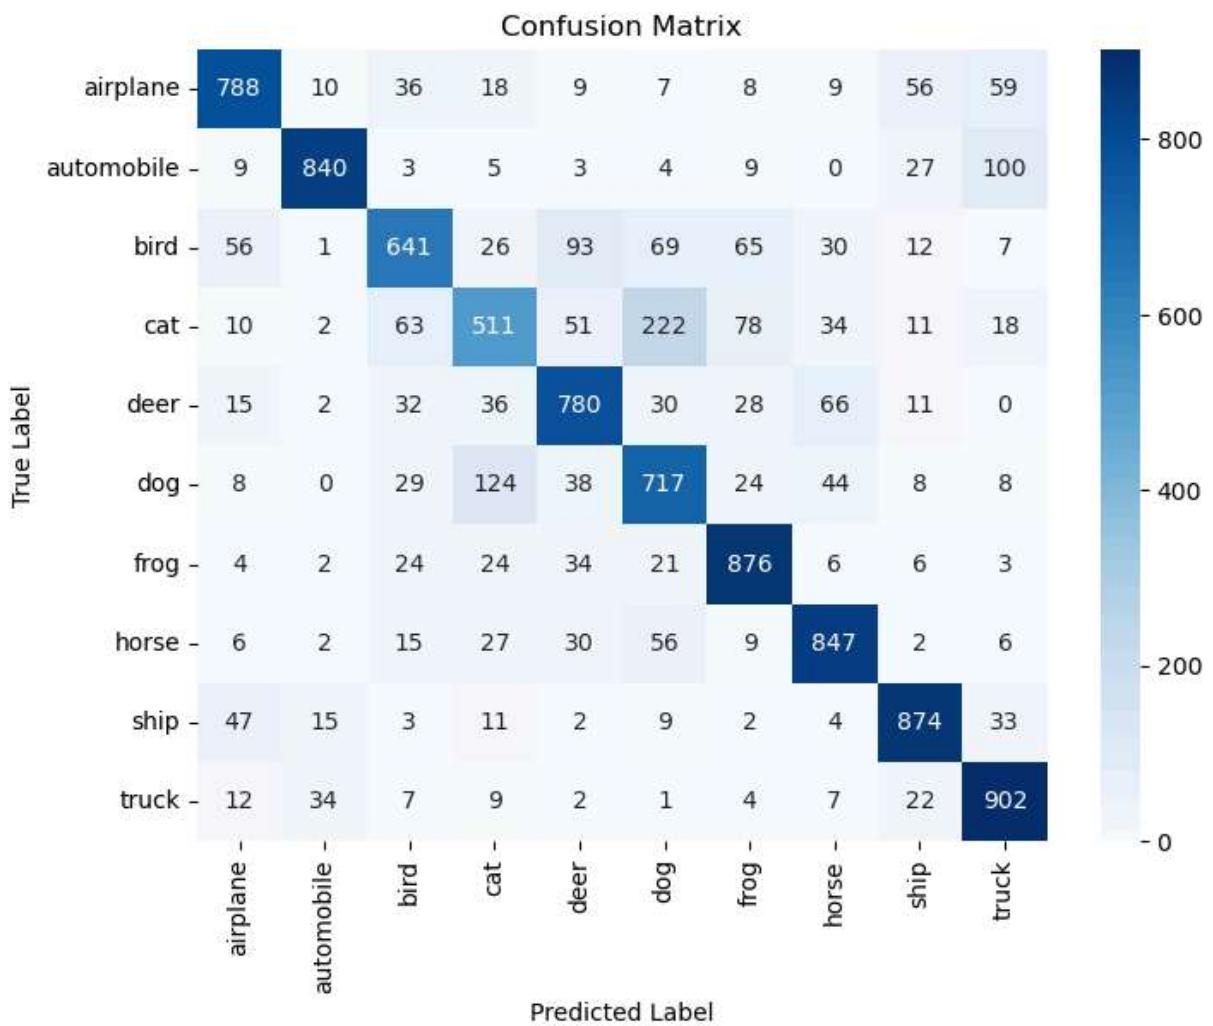
```
In [67]: y_pred = model.predict(x_test)
y_pred_classes = np.argmax(y_pred, axis=1)
y_test_classes = y_test.flatten()
```

```
313/313 ————— 6s 18ms/step
```

```
In [69]: cm = confusion_matrix(y_test_classes, y_pred_classes)
```

```
In [71]: class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                    'dog', 'frog', 'horse', 'ship', 'truck']
```

```
In [73]: plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=class_names, yticklabels=class_names)
plt.xlabel('Predicted Label')
plt.ylabel('True Label')
plt.title('Confusion Matrix')
plt.show()
```



```
In [75]: print("Classification Report:\n", classification_report(y_test_classes, y_pred_classes))
```

	precision	recall	f1-score	support
airplane	0.83	0.79	0.81	1000
automobile	0.93	0.84	0.88	1000
bird	0.75	0.64	0.69	1000
cat	0.65	0.51	0.57	1000
deer	0.75	0.78	0.76	1000
dog	0.63	0.72	0.67	1000
frog	0.79	0.88	0.83	1000
horse	0.81	0.85	0.83	1000
ship	0.85	0.87	0.86	1000
truck	0.79	0.90	0.84	1000
accuracy			0.78	10000
macro avg	0.78	0.78	0.78	10000
weighted avg	0.78	0.78	0.78	10000

Insights from Classification Report:

Accuracy: 78%

Examples of Correctly and Incorrectly Classified Images

Precision & Recall: The model struggles with Cats and Birds, as their F1-scores are low (~0.57-0.69), meaning they are often misclassified.

Strong Performance: The model correctly identifies Ships, Automobiles, and Horses with high precision and recall (~0.85+), showing confidence in these categories.

Dogs & Deer: The model shows imbalanced recall and precision, meaning it sometimes misidentifies them or confuses them with similar classes.

Overall: The model performs fairly well but struggles with certain categories. Some tuning or better data augmentation might help improve its performance.

```
In [79]: from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical
```

Here I am loading the CIFAR-10 dataset, preprocessing it, training the model, and then evaluating its performance on test data.

```
In [81]: (x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train, y_train, x_val, y_val, x_test, y_test = preprocess_data(x_train, y_train,
model = build_model()
history = train_model(model, x_train, y_train, x_val, y_val)
evaluate_model(model, x_test, y_test)
```

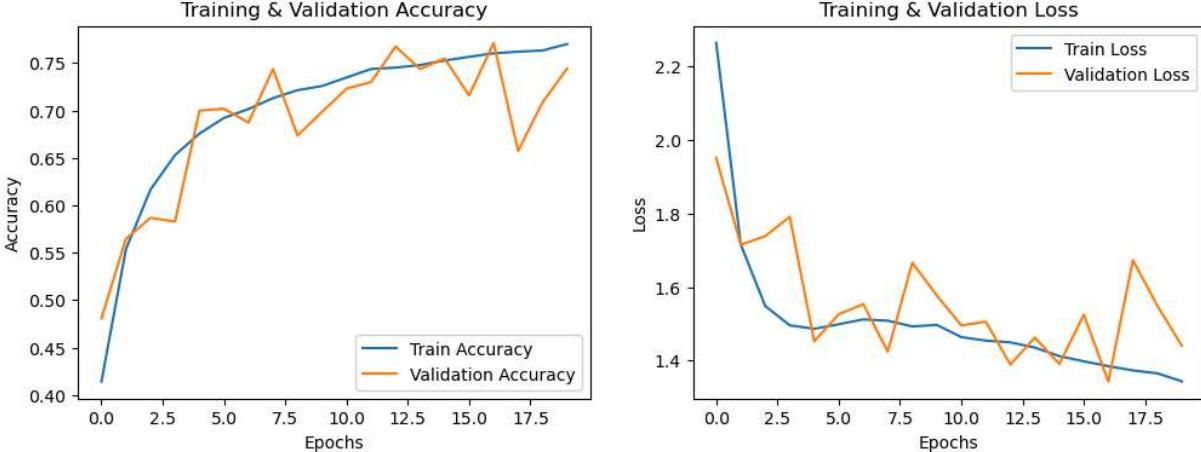
```
C:\Users\msgme\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
Epoch 1/20
704/704 139s 190ms/step - accuracy: 0.3384 - loss: 2.8308 - val
_accuracy: 0.4808 - val_loss: 1.9515
Epoch 2/20
704/704 136s 193ms/step - accuracy: 0.5389 - loss: 1.7730 - val
_accuracy: 0.5646 - val_loss: 1.7150
Epoch 3/20
704/704 138s 196ms/step - accuracy: 0.6063 - loss: 1.5621 - val
_accuracy: 0.5868 - val_loss: 1.7382
Epoch 4/20
704/704 142s 201ms/step - accuracy: 0.6493 - loss: 1.4935 - val
_accuracy: 0.5828 - val_loss: 1.7911
Epoch 5/20
704/704 140s 199ms/step - accuracy: 0.6809 - loss: 1.4592 - val
_accuracy: 0.6998 - val_loss: 1.4525
Epoch 6/20
704/704 142s 201ms/step - accuracy: 0.6930 - loss: 1.4859 - val
_accuracy: 0.7018 - val_loss: 1.5261
Epoch 7/20
704/704 142s 202ms/step - accuracy: 0.7047 - loss: 1.4889 - val
_accuracy: 0.6872 - val_loss: 1.5535
Epoch 8/20
704/704 141s 200ms/step - accuracy: 0.7174 - loss: 1.4857 - val
_accuracy: 0.7436 - val_loss: 1.4245
Epoch 9/20
704/704 149s 212ms/step - accuracy: 0.7280 - loss: 1.4703 - val
_accuracy: 0.6736 - val_loss: 1.6661
Epoch 10/20
704/704 169s 240ms/step - accuracy: 0.7306 - loss: 1.4713 - val
_accuracy: 0.6984 - val_loss: 1.5773
Epoch 11/20
704/704 169s 239ms/step - accuracy: 0.7419 - loss: 1.4423 - val
_accuracy: 0.7230 - val_loss: 1.4958
Epoch 12/20
704/704 168s 239ms/step - accuracy: 0.7438 - loss: 1.4500 - val
_accuracy: 0.7298 - val_loss: 1.5060
Epoch 13/20
704/704 169s 240ms/step - accuracy: 0.7525 - loss: 1.4232 - val
_accuracy: 0.7674 - val_loss: 1.3891
Epoch 14/20
704/704 168s 239ms/step - accuracy: 0.7543 - loss: 1.4192 - val
_accuracy: 0.7436 - val_loss: 1.4628
Epoch 15/20
704/704 169s 240ms/step - accuracy: 0.7603 - loss: 1.3947 - val
_accuracy: 0.7548 - val_loss: 1.3909
Epoch 16/20
704/704 170s 241ms/step - accuracy: 0.7605 - loss: 1.3786 - val
_accuracy: 0.7158 - val_loss: 1.5251
Epoch 17/20
704/704 171s 243ms/step - accuracy: 0.7643 - loss: 1.3714 - val
_accuracy: 0.7712 - val_loss: 1.3430
Epoch 18/20
704/704 172s 244ms/step - accuracy: 0.7683 - loss: 1.3534 - val
_accuracy: 0.6572 - val_loss: 1.6730
Epoch 19/20
704/704 171s 242ms/step - accuracy: 0.7674 - loss: 1.3517 - val
```

```

_accuracy: 0.7088 - val_loss: 1.5491
Epoch 20/20
704/704 ————— 169s 241ms/step - accuracy: 0.7745 - loss: 1.3260 - val
_accuracy: 0.7442 - val_loss: 1.4402

```



```

313/313 ————— 10s 32ms/step - accuracy: 0.7715 - loss: 1.3624
Test Accuracy: 0.7670000195503235

```

Model learned well and got 77.45% accuracy on training data, 74.42% on validation data, and 76.7% on test data.

This means it's doing a good job but slightly overfits. The training loss kept decreasing, but the validation loss

went up and down, showing the model struggled a bit with new data. To improve, The model can improve by using try

Dropout, L2 regularization, or data augmentation to make it more stable. Also, Improving settings like the learning

rate or optimizer might boost accuracy!

Task 4: Experimentation with Model Improvements.

Comparing Different Optimizers for Model Performance

```
In [86]: def experiment_with_optimizers():
    optimizers_list = ['adam', 'sgd', 'rmsprop']
    results={}
    for opt in optimizers_list:
        print(f"\nTraining with {opt} optimizer:")
        model = build_model()
        model.compile(optimizer=opt, loss='sparse_categorical_crossentropy', metrics=['accuracy'])
        history = train_model(model, x_train, y_train, x_val, y_val)
        test_loss, test_acc = model.evaluate(x_test, y_test)
        results[opt]=test_acc
    print("\nPerformance Comparison:")
    for opt, acc in results.items():
        print(f"{opt}: Test Accuracy = {acc:.4f}")
    return results
```

Creating a function to tests three different optimizers (Adam, SGD, and RMSprop) to see which one gives the best accuracy.

It will train a model with each optimizer, evaluate its performance, and compare the test accuracy of all three. It will

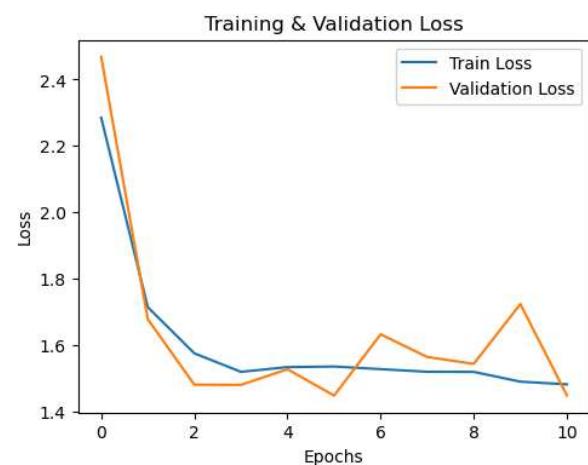
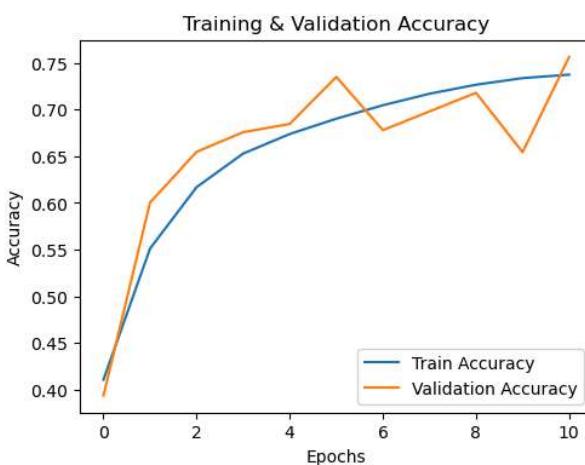
help in choosing the best optimizer for better model performance.

Comparing Optimizers Adam vs. SGD vs. RMSprop to Find the Best One!

```
In [91]: results = experiment_with_optimizers()
```

Training with adam optimizer:

```
Epoch 1/20  
704/704 183s 244ms/step - accuracy: 0.3334 - loss: 2.9400 - val_accuracy: 0.3938 - val_loss: 2.4682  
Epoch 2/20  
704/704 171s 243ms/step - accuracy: 0.5285 - loss: 1.7781 - val_accuracy: 0.6004 - val_loss: 1.6773  
Epoch 3/20  
704/704 172s 244ms/step - accuracy: 0.6087 - loss: 1.5871 - val_accuracy: 0.6548 - val_loss: 1.4800  
Epoch 4/20  
704/704 172s 244ms/step - accuracy: 0.6507 - loss: 1.5186 - val_accuracy: 0.6758 - val_loss: 1.4793  
Epoch 5/20  
704/704 171s 243ms/step - accuracy: 0.6721 - loss: 1.5132 - val_accuracy: 0.6846 - val_loss: 1.5266  
Epoch 6/20  
704/704 171s 242ms/step - accuracy: 0.6945 - loss: 1.5114 - val_accuracy: 0.7350 - val_loss: 1.4470  
Epoch 7/20  
704/704 170s 242ms/step - accuracy: 0.7101 - loss: 1.4950 - val_accuracy: 0.6780 - val_loss: 1.6318  
Epoch 8/20  
704/704 171s 242ms/step - accuracy: 0.7206 - loss: 1.5034 - val_accuracy: 0.6980 - val_loss: 1.5634  
Epoch 9/20  
704/704 171s 243ms/step - accuracy: 0.7271 - loss: 1.5082 - val_accuracy: 0.7180 - val_loss: 1.5428  
Epoch 10/20  
704/704 171s 243ms/step - accuracy: 0.7377 - loss: 1.4793 - val_accuracy: 0.6544 - val_loss: 1.7231  
Epoch 11/20  
704/704 171s 243ms/step - accuracy: 0.7433 - loss: 1.4600 - val_accuracy: 0.7564 - val_loss: 1.4474
```



313/313 ━━━━━━ 10s 31ms/step - accuracy: 0.7284 - loss: 1.4646

Training with sgd optimizer:

Epoch 1/20

704/704 ━━━━━━ 157s 218ms/step - accuracy: 0.2879 - loss: 2.8484 - val_accuracy: 0.4628 - val_loss: 2.1239

Epoch 2/20

704/704 ━━━━━━ 126s 178ms/step - accuracy: 0.4413 - loss: 2.1069 - val_accuracy: 0.5200 - val_loss: 1.9368

Epoch 3/20

704/704 ━━━━━━ 128s 182ms/step - accuracy: 0.4894 - loss: 1.9705 - val_accuracy: 0.5134 - val_loss: 2.0517

Epoch 4/20

704/704 ━━━━━━ 128s 182ms/step - accuracy: 0.5190 - loss: 1.8908 - val_accuracy: 0.4932 - val_loss: 1.9756

Epoch 5/20

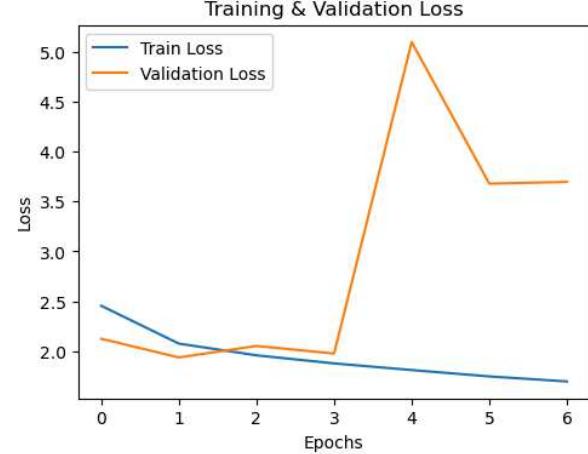
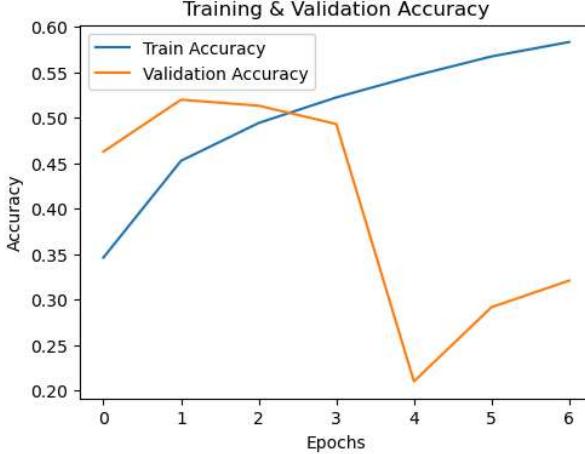
704/704 ━━━━━━ 129s 184ms/step - accuracy: 0.5399 - loss: 1.8217 - val_accuracy: 0.2098 - val_loss: 5.1013

Epoch 6/20

704/704 ━━━━━━ 128s 182ms/step - accuracy: 0.5646 - loss: 1.7521 - val_accuracy: 0.2916 - val_loss: 3.6802

Epoch 7/20

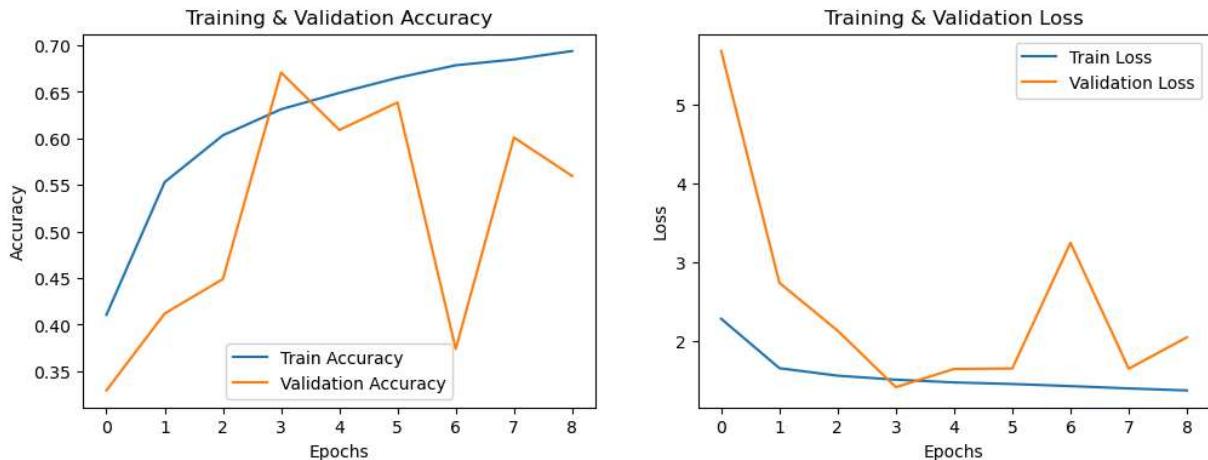
704/704 ━━━━━━ 127s 181ms/step - accuracy: 0.5800 - loss: 1.7066 - val_accuracy: 0.3208 - val_loss: 3.6975



313/313 ————— 7s 22ms/step - accuracy: 0.5118 - loss: 1.9487

Training with rmsprop optimizer:

Epoch 1/20
704/704 ————— 137s 189ms/step - accuracy: 0.3354 - loss: 3.2163 - val_accuracy: 0.3294 - val_loss: 5.6820
 Epoch 2/20
704/704 ————— 136s 193ms/step - accuracy: 0.5366 - loss: 1.6880 - val_accuracy: 0.4120 - val_loss: 2.7398
 Epoch 3/20
704/704 ————— 136s 194ms/step - accuracy: 0.5965 - loss: 1.5672 - val_accuracy: 0.4490 - val_loss: 2.1329
 Epoch 4/20
704/704 ————— 137s 194ms/step - accuracy: 0.6286 - loss: 1.5126 - val_accuracy: 0.6708 - val_loss: 1.4158
 Epoch 5/20
704/704 ————— 138s 197ms/step - accuracy: 0.6448 - loss: 1.4682 - val_accuracy: 0.6090 - val_loss: 1.6472
 Epoch 6/20
704/704 ————— 139s 197ms/step - accuracy: 0.6654 - loss: 1.4482 - val_accuracy: 0.6386 - val_loss: 1.6529
 Epoch 7/20
704/704 ————— 138s 196ms/step - accuracy: 0.6806 - loss: 1.4246 - val_accuracy: 0.3740 - val_loss: 3.2477
 Epoch 8/20
704/704 ————— 141s 200ms/step - accuracy: 0.6874 - loss: 1.3863 - val_accuracy: 0.6010 - val_loss: 1.6501
 Epoch 9/20
704/704 ————— 139s 197ms/step - accuracy: 0.6920 - loss: 1.3703 - val_accuracy: 0.5596 - val_loss: 2.0491



313/313 ————— 8s 25ms/step - accuracy: 0.6700 - loss: 1.4248

Performance Comparison:

adam: Test Accuracy = 0.7271
 sgd: Test Accuracy = 0.5046
 rmsprop: Test Accuracy = 0.6678

Displaying Test Accuracy for Each Optimizer!

```
In [99]: for opt, acc in results.items():
    print(f"{opt}: Test Accuracy = {acc:.4f}")
```

```
adam: Test Accuracy = 0.7271
sgd: Test Accuracy = 0.5046
rmsprop: Test Accuracy = 0.6678
```

I tested three optimizers—Adam, SGD, and RMSprop—to find the best one. Adam performed the best with 72.71% accuracy,

while RMSprop with 66.78%, and SGD struggled at 50.46%. This means Adam helped the model learn better, but

SGD didn't perform well. The result indicates that choosing the right optimizer is important for getting good accuracy!

```
In [ ]:
```